

Part I_Proprosper Loan Exploration

January 10, 2023

1 Prosper Loan Data Exploration

1.1 by Emma Hungrige

1.2 Preliminary Wrangling

This data set contains information on peer to peer loans facilitated by the credit company, Prosper.

```
In [74]: # import all packages and set plots to be embedded inline
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings("ignore")

%matplotlib inline
```

```
%config InlineBackend.figure_format = "retina"
```

```
In [75]: df = pd.read_csv("prosperLoanData.csv")
df.head()
```

```
Out[75]:
```

	ListingKey	ListingNumber	ListingCreationDate	
0	1021339766868145413AB3B	193129	2007-08-26 19:09:29.263000000	
1	10273602499503308B223C1	1209647	2014-02-27 08:28:07.900000000	
2	0EE9337825851032864889A	81716	2007-01-05 15:00:47.090000000	
3	0EF5356002482715299901A	658116	2012-10-22 11:02:35.010000000	
4	0F023589499656230C5E3E2	909464	2013-09-14 18:38:39.097000000	

	CreditGrade	Term	LoanStatus	ClosedDate	BorrowerAPR	
0	C	36	Completed	2009-08-14 00:00:00	0.16516	
1	NaN	36	Current	NaN	0.12016	
2	HR	36	Completed	2009-12-17 00:00:00	0.28269	
3	NaN	36	Current	NaN	0.12528	
4	NaN	36	Current	NaN	0.24614	

	BorrowerRate	LenderYield	...	LP_ServiceFees	LP_CollectionFees	\
0	0.1580	0.1380	...	-133.18	0.0	
1	0.0920	0.0820	...	0.00	0.0	
2	0.2750	0.2400	...	-24.20	0.0	
3	0.0974	0.0874	...	-108.01	0.0	
4	0.2085	0.1985	...	-60.27	0.0	

	LP_GrossPrincipalLoss	LP_NetPrincipalLoss	LP_NonPrincipalRecoverypayments	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	

	PercentFunded	Recommendations	InvestmentFromFriendsCount	\
0	1.0	0	0	
1	1.0	0	0	
2	1.0	0	0	
3	1.0	0	0	
4	1.0	0	0	

	InvestmentFromFriendsAmount	Investors
0	0.0	258
1	0.0	1
2	0.0	41
3	0.0	158
4	0.0	20

[5 rows x 81 columns]

In [76]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 113937 entries, 0 to 113936
Data columns (total 81 columns):
```

#	Column	Non-Null Count	Dtype
0	ListingKey	113937 non-null	object
1	ListingNumber	113937 non-null	int64
2	ListingCreationDate	113937 non-null	object
3	CreditGrade	28953 non-null	object
4	Term	113937 non-null	int64
5	LoanStatus	113937 non-null	object
6	ClosedDate	55089 non-null	object
7	BorrowerAPR	113912 non-null	float64
8	BorrowerRate	113937 non-null	float64
9	LenderYield	113937 non-null	float64
10	EstimatedEffectiveYield	84853 non-null	float64

11	EstimatedLoss	84853 non-null	float64
12	EstimatedReturn	84853 non-null	float64
13	ProsperRating (numeric)	84853 non-null	float64
14	ProsperRating (Alpha)	84853 non-null	object
15	ProsperScore	84853 non-null	float64
16	ListingCategory (numeric)	113937 non-null	int64
17	BorrowerState	108422 non-null	object
18	Occupation	110349 non-null	object
19	EmploymentStatus	111682 non-null	object
20	EmploymentStatusDuration	106312 non-null	float64
21	IsBorrowerHomeowner	113937 non-null	bool
22	CurrentlyInGroup	113937 non-null	bool
23	GroupKey	13341 non-null	object
24	DateCreditPulled	113937 non-null	object
25	CreditScoreRangeLower	113346 non-null	float64
26	CreditScoreRangeUpper	113346 non-null	float64
27	FirstRecordedCreditLine	113240 non-null	object
28	CurrentCreditLines	106333 non-null	float64
29	OpenCreditLines	106333 non-null	float64
30	TotalCreditLinespast7years	113240 non-null	float64
31	OpenRevolvingAccounts	113937 non-null	int64
32	OpenRevolvingMonthlyPayment	113937 non-null	float64
33	InquiriesLast6Months	113240 non-null	float64
34	TotalInquiries	112778 non-null	float64
35	CurrentDelinquencies	113240 non-null	float64
36	AmountDelinquent	106315 non-null	float64
37	DelinquenciesLast7Years	112947 non-null	float64
38	PublicRecordsLast10Years	113240 non-null	float64
39	PublicRecordsLast12Months	106333 non-null	float64
40	RevolvingCreditBalance	106333 non-null	float64
41	BankcardUtilization	106333 non-null	float64
42	AvailableBankcardCredit	106393 non-null	float64
43	TotalTrades	106393 non-null	float64
44	TradesNeverDelinquent (percentage)	106393 non-null	float64
45	TradesOpenedLast6Months	106393 non-null	float64
46	DebtToIncomeRatio	105383 non-null	float64
47	IncomeRange	113937 non-null	object
48	IncomeVerifiable	113937 non-null	bool
49	StatedMonthlyIncome	113937 non-null	float64
50	LoanKey	113937 non-null	object
51	TotalProsperLoans	22085 non-null	float64
52	TotalProsperPaymentsBilled	22085 non-null	float64
53	OnTimeProsperPayments	22085 non-null	float64
54	ProsperPaymentsLessThanOneMonthLate	22085 non-null	float64
55	ProsperPaymentsOneMonthPlusLate	22085 non-null	float64
56	ProsperPrincipalBorrowed	22085 non-null	float64
57	ProsperPrincipalOutstanding	22085 non-null	float64
58	ScorexChangeAtTimeOfListing	18928 non-null	float64

```

59 LoanCurrentDaysDelinquent      113937 non-null int64
60 LoanFirstDefaultedCycleNumber  16952 non-null float64
61 LoanMonthsSinceOrigination     113937 non-null int64
62 LoanNumber                     113937 non-null int64
63 LoanOriginalAmount             113937 non-null int64
64 LoanOriginationDate            113937 non-null object
65 LoanOriginationQuarter         113937 non-null object
66 MemberKey                     113937 non-null object
67 MonthlyLoanPayment             113937 non-null float64
68 LP_CustomerPayments            113937 non-null float64
69 LP_CustomerPrincipalPayments   113937 non-null float64
70 LP_InterestandFees            113937 non-null float64
71 LP_ServiceFees                 113937 non-null float64
72 LP_CollectionFees             113937 non-null float64
73 LP_GrossPrincipalLoss          113937 non-null float64
74 LP_NetPrincipalLoss           113937 non-null float64
75 LP_NonPrincipalRecoverypayments 113937 non-null float64
76 PercentFunded                 113937 non-null float64
77 Recommendations                113937 non-null int64
78 InvestmentFromFriendsCount     113937 non-null int64
79 InvestmentFromFriendsAmount    113937 non-null float64
80 Investors                     113937 non-null int64

```

dtypes: bool(3), float64(50), int64(11), object(17)

memory usage: 68.1+ MB

In [77]: df.describe()

```

Out[77]:

```

	ListingNumber	Term	BorrowerAPR	BorrowerRate	\
count	1.139370e+05	113937.000000	113912.000000	113937.000000	
mean	6.278857e+05	40.830248	0.218828	0.192764	
std	3.280762e+05	10.436212	0.080364	0.074818	
min	4.000000e+00	12.000000	0.006530	0.000000	
25%	4.009190e+05	36.000000	0.156290	0.134000	
50%	6.005540e+05	36.000000	0.209760	0.184000	
75%	8.926340e+05	36.000000	0.283810	0.250000	
max	1.255725e+06	60.000000	0.512290	0.497500	

	LenderYield	EstimatedEffectiveYield	EstimatedLoss	EstimatedReturn	\
count	113937.000000	84853.000000	84853.000000	84853.000000	
mean	0.182701	0.168661	0.080306	0.096068	
std	0.074516	0.068467	0.046764	0.030403	
min	-0.010000	-0.182700	0.004900	-0.182700	
25%	0.124200	0.115670	0.042400	0.074080	
50%	0.173000	0.161500	0.072400	0.091700	
75%	0.240000	0.224300	0.112000	0.116600	
max	0.492500	0.319900	0.366000	0.283700	

	ProsperRating (numeric)	ProsperScore	...	LP_ServiceFees \
count	84853.000000	84853.000000	...	113937.000000
mean	4.072243	5.950067	...	-54.725641
std	1.673227	2.376501	...	60.675425
min	1.000000	1.000000	...	-664.870000
25%	3.000000	4.000000	...	-73.180000
50%	4.000000	6.000000	...	-34.440000
75%	5.000000	8.000000	...	-13.920000
max	7.000000	11.000000	...	32.060000

	LP_CollectionFees	LP_GrossPrincipalLoss	LP_NetPrincipalLoss \
count	113937.000000	113937.000000	113937.000000
mean	-14.242698	700.446342	681.420499
std	109.232758	2388.513831	2357.167068
min	-9274.750000	-94.200000	-954.550000
25%	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000
max	0.000000	25000.000000	25000.000000

	LP_NonPrincipalRecoverypayments	PercentFunded	Recommendations \
count	113937.000000	113937.000000	113937.000000
mean	25.142686	0.998584	0.048027
std	275.657937	0.017919	0.332353
min	0.000000	0.700000	0.000000
25%	0.000000	1.000000	0.000000
50%	0.000000	1.000000	0.000000
75%	0.000000	1.000000	0.000000
max	21117.900000	1.012500	39.000000

	InvestmentFromFriendsCount	InvestmentFromFriendsAmount	Investors
count	113937.000000	113937.000000	113937.000000
mean	0.023460	16.550751	80.475228
std	0.232412	294.545422	103.239020
min	0.000000	0.000000	1.000000
25%	0.000000	0.000000	2.000000
50%	0.000000	0.000000	44.000000
75%	0.000000	0.000000	115.000000
max	33.000000	25000.000000	1189.000000

[8 rows x 61 columns]

In [78]: df.sample(10)

Out[78]:	ListingKey	ListingNumber	ListingCreationDate \
2245	041F3430145017632AA8BA0	390137	2008-08-29 05:07:43.680000000
75013	40F535908073254938858BC	912400	2013-09-19 13:39:32.590000000
46273	8DOC3530238955853B93AA2	534978	2011-10-23 14:07:59.283000000

102988	35113586293088828096405	859552	2013-08-03 10:46:43.993000000
64366	91083586611460835DA94F3	862308	2013-08-07 03:22:49.683000000
33939	B8663596881629915202A61	1100583	2013-12-20 11:02:19.417000000
77865	9DA635461878864198A8F59	588062	2012-05-10 12:40:01.903000000
1379	7B22358909824432888A960	913677	2013-09-20 09:47:04.787000000
58245	BC6A3557201311989982118	634327	2012-09-05 12:37:31.397000000
27156	01793403613525474C385A5	228220	2007-11-08 04:22:35.990000000

	CreditGrade	Term	LoanStatus	ClosedDate \
2245	B	36	Completed	2011-09-09 00:00:00
75013	NaN	36	Current	NaN
46273	NaN	36	Completed	2012-06-15 00:00:00
102988	NaN	60	Current	NaN
64366	NaN	60	Current	NaN
33939	NaN	60	Current	NaN
77865	NaN	36	Past Due (1-15 days)	NaN
1379	NaN	60	Current	NaN
58245	NaN	36	Completed	2013-10-10 00:00:00
27156	D	36	Completed	2008-05-16 00:00:00

	BorrowerAPR	BorrowerRate	LenderYield	...	LP_ServiceFees \
2245	0.17193	0.1575	0.1475	...	-57.94
75013	0.14409	0.1159	0.1059	...	-35.93
46273	0.20200	0.1799	0.1699	...	-14.74
102988	0.27637	0.2506	0.2406	...	-19.73
64366	0.13942	0.1169	0.1069	...	-73.25
33939	0.20040	0.1765	0.1665	...	-16.90
77865	0.35797	0.3177	0.3077	...	-47.68
1379	0.21115	0.1870	0.1770	...	-41.04
58245	0.27060	0.2324	0.2224	...	-17.18
27156	0.22868	0.2100	0.2000	...	-4.78

	LP_CollectionFees	LP_GrossPrincipalLoss	LP_NetPrincipalLoss \
2245	0.00	0.0	0.0
75013	0.00	0.0	0.0
46273	0.00	0.0	0.0
102988	0.00	0.0	0.0
64366	0.00	0.0	0.0
33939	0.00	0.0	0.0
77865	-147.65	0.0	0.0
1379	0.00	0.0	0.0
58245	0.00	0.0	0.0
27156	0.00	0.0	0.0

	LP_NonPrincipalRecoverypayments	PercentFunded	Recommendations \
2245	0.0	1.0	0
75013	0.0	1.0	0
46273	0.0	1.0	0

102988	0.0	1.0	0
64366	0.0	1.0	0
33939	0.0	1.0	0
77865	0.0	1.0	0
1379	0.0	1.0	0
58245	0.0	1.0	0
27156	0.0	1.0	0

	InvestmentFromFriendsCount	InvestmentFromFriendsAmount	Investors
2245	0	0.0	108
75013	0	0.0	90
46273	0	0.0	20
102988	0	0.0	2
64366	0	0.0	274
33939	0	0.0	1
77865	0	0.0	13
1379	0	0.0	1
58245	0	0.0	29
27156	0	0.0	14

[10 rows x 81 columns]

1.2.1 What is/are the main feature(s) of interest in your dataset?

We're most interested in figuring out which borrower features are best for predicting the highest rate of return and what factors will have the highest impact on the chances of default. We also want to investigate how closely the estimated loan performance matches the actual loan performance.

1.2.2 What features in the dataset do you think will help support your investigation into your feature(s) of interest?

There are over 81 columns to use in this dataset, but we believe the following columns will have the largest effect on the performance of the loans:

1. **IncomeRange** - The income range of the borrower at the time the listing was created
2. **DebtToIncomeRatio** - The debt to income ratio of the borrower at the time the credit profile was pulled. This value is Null if the debt to income ratio is not available. This value is capped at 10.01 (any debt to income ratio larger than 1000% will be returned at 1001%).
3. **CreditGrade** - The Credit rating that was assigned at the time the listing went live. Applicable for listings pre-2009 period and will only be populated for those listings.
4. **ProsperRating (Alpha)** - The Prosper Rating assigned at the time the listing was created between AA - HR. Applicable for loans originated after July 2009.

The performance of the loan can be determined by the following features:

1. **EstimatedReturn** - The estimated return assigned to the listing at the time it was created. Estimated return is the difference between the Estimated Effective Yield and the Estimated

Loss Rate. Applicable for loans originated after July 2009. This is on an annual percentage rate (APR).

2. **ActualReturn** - To see how well the loan performed verses the EstimatedReturn. We will calculate this feature based on how Prosper calculates their Annualized Net Returns.

1.3 Univariate Exploration

1.3.1 Estimated Return

```
In [79]: # Setting color
         base_color = sns.color_palette()[0]
```

```
In [80]: df.EstimatedReturn.describe()
```

```
Out[80]: count      84853.000000
         mean         0.096068
         std         0.030403
         min        -0.182700
         25%         0.074080
         50%         0.091700
         75%         0.116600
         max         0.283700
         Name: EstimatedReturn, dtype: float64
```

There appears to be quite a few empty cells as the total number of rows is over 10k.


The range of estimated returns is -18% to 28%.

I expect to see a heavily right skewed distribution plot based on the quartile distribution.

Let's drop the null rows and create a histogram plot.

```
In [81]: #Drop null rows without an estimated return
         df.EstimatedReturn.dropna(axis = 0, inplace = True)
```


```
In [82]: #Histogram plot
         plt.figure(figsize = [8, 5])
         sns.distplot(df.EstimatedReturn, kde = False, bins = 200, color = base_color)
         plt.xlabel("Estimated Return")
         plt.ylabel("Count");
```

It appears most of our loans have an estimated return between 0% and 20%. An additional histogram shall be created to zoom in a little further.

In [83]: *#Zoom in on histogram*

```
plt.figure(figsize = [8,5])
sns.distplot(df.EstimatedReturn, kde = False, bins = 200, color = base_color)
plt.xlim(0, .2)
plt.xlabel("Estimated Return")
plt.ylabel("Count")
plt.title("Distribution of Estimated Returns");
```



It appears like the assumption that this data would be right skewed is correct.

It also looks like there are a few peaks in some standard values like 5%, 7.5%, and 15%. Most interestingly, the largest bin falls right around 12.5%.

1.3.2 Actual Returns

This section appears to be quite nuanced and required some engineering. We're going to calculate the actual returns by using the remaining loan performance (LP) variables.

NOTE: To simplify out calculation for exploratory purposes, we will only divide by the original loan amount. Otherwise, we'll run into other categorical issues if we try to determine the current rate of return for current loans vs completed loans. This section may be revisited later to refine the actual returns.


```
In [84]: #Simplified actual return formula
df["ActualReturns"] = df.LP_CustomerPayments - df.LoanOriginalAmount - (df.LP_ServiceFee
df["ActualReturns"] = 1 + (df.ActualReturns / df.LoanOriginalAmount)
```

```
In [85]: df.ActualReturns.describe()
```

```
Out[85]: count      113937.000000
         mean         0.531620
         std         0.621164
         min        -1.000650
         25%         0.074989
         50%         0.415864
         75%         1.119834
```

```
max                2.602880
Name: ActualReturns, dtype: float64
```

```
In [86]: #Plot distribution
plt.figure(figsize = [8,5])
sns.distplot(df.ActualReturns, kde = False, bins = 200, color = base_color)
plt.xlim(-1.1, 1.8)
plt.xlabel("Actual Returns")
plt.ylabel("Count");
```



It's interesting to see there are some loans where it appears that the borrower never made a payment which gives us the minimum of losing 100% of the investment. On the other end of the spectrum, it looks like most loans drop off after about 170% of the original loan amount.

As mentioned previously, to simplify the visualization, this calculation is not annualized based on the length of the loans, so this is not quite an apples to apples comparison with the estimated return which is annualized.

Since there appears to be different modes to the distribution, let's see how the actual returns differ by the loan status:

```
In [87]: #Multiplot for loan status
```

```
variables = ["Current", "Completed", "Chargedoff", "Defaulted"]
```

```
plt.figure(figsize = [16, 10])
```

```
# loop through variables list
```

```

for i in range(len(variables)):
    plt.subplot(2, 2, i+1)

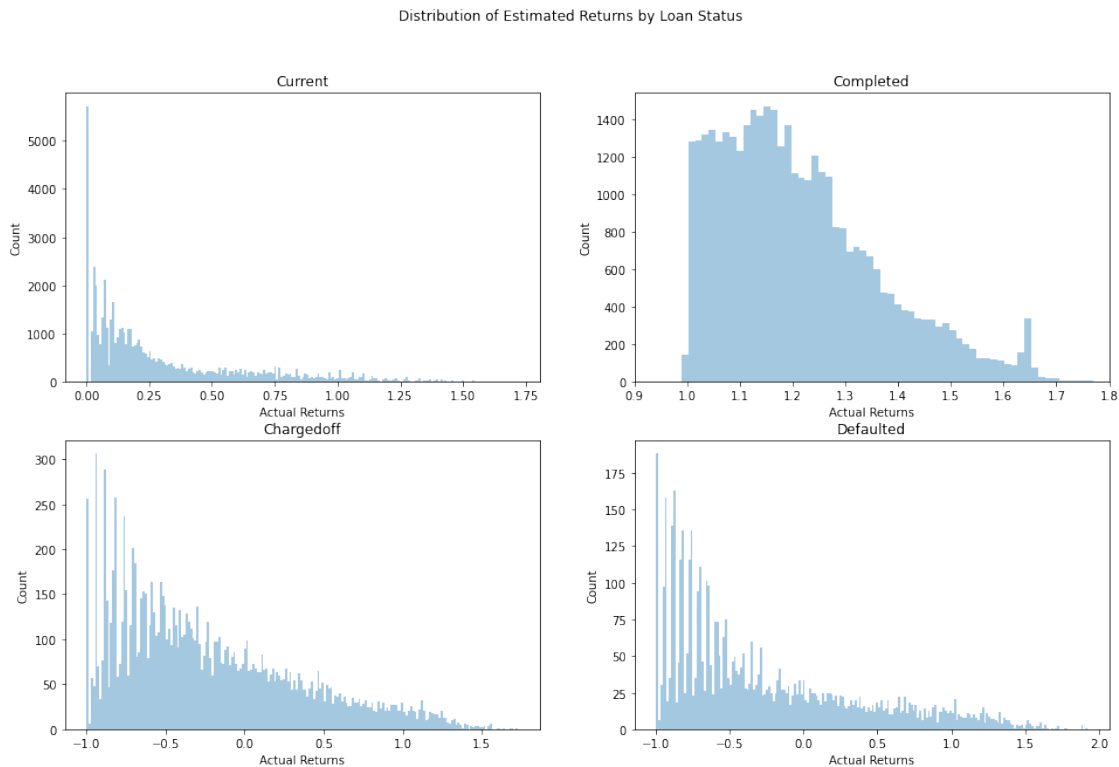
    sns.distplot(df.query('LoanStatus == "{}"'.format(variables[i])).ActualReturns,
                  kde = False, bins = 200, color = base_color)

    plt.xlabel('Actual Returns')
    plt.ylabel('Count')
    plt.title(variables[i])

plt.subplot(2, 2, 2)
plt.xlim(.9, 1.8)

plt.suptitle('Distribution of Estimated Returns by Loan Status');

```



For the current loans, there appears to be a large number of loans that have an actual return value of 0. This could mean that "Current" also takes into account all the loans that have not been fully funded yet. We would have to explore this further to extract loans that have been fully funded.

Since our actual returns feature is not annualized, we see that for the loan terms of 3-5 years, an investor could see a total return of 170%.

For loans that have defaulted or have been chargedoff, it makes sense that a majority of them show negative returns up to 100%. Meaning that the borrower could have never made a payment

or the service/collection fees resulted in the investor losing their entire principal. Of course this is the worst case scenario.

1.3.3 Prosper Credit Grade

Now that we have a high level overview of what the returns on the platform looks like, let's investigate if there is a similar distribution for the borrower's prosper/credit grade.

```
In [88]: df['ProsperRating (Alpha)'].describe(), df['ProsperRating (Alpha)'].value_counts()
```

```
Out[88]: (count      84853
         unique        7
         top          C
         freq      18345
         Name: ProsperRating (Alpha), dtype: object,
         C      18345
         B      15581
         A      14551
         D      14274
         E       9795
         HR       6935
         AA       5372
         Name: ProsperRating (Alpha), dtype: int64)
```

```
In [89]: df['CreditGrade'].describe(), df['CreditGrade'].value_counts()
```

```
Out[89]: (count      28953
         unique        8
         top          C
         freq      5649
         Name: CreditGrade, dtype: object,
         C      5649
         D      5153
         B      4389
         AA      3509
         HR      3508
         A      3315
         E      3289
         NC       141
         Name: CreditGrade, dtype: int64)
```

From the feature descriptions, it looks like Prosper transitioned to a "Prosper Rating" after 2009 to differentiate themselves from the industry standard "Credit Grade".

For exploratory purposes, we'll replace the Credit Grade of No Credit (NC) with High Risk (HR) and combine this column with the Prosper Rating.

In the future we may want to look at estimated returns of the pre-2009 "Credit Grade" to see how it compares with the post-2009 "Prosper Rating" to validate if they truly are equivalent or if Prosper skewed their credit rating in some way.

```

In [90]: #Replace NC with HR
df["CreditGrade"].replace(["NC", "HR"], "E", inplace = True)

#Fill NA with empty strings
df["ProsperRating (Alpha)"].fillna("", inplace = True)
df["CreditGrade"].fillna("", inplace = True)

#Combine credit ratings
df["ProsperRatingCombined"] = df["ProsperRating (Alpha)"] + df["CreditGrade"]

#Drop rows with no credit ratings
df = df.query(' ProsperRatingCombined != ""')

#Check values
df["ProsperRatingCombined"].value_counts()

```

```

Out[90]: C      23994
        B      19970
        D      19427
        A      17866
        E      16733
        AA       8881
        HR       6935
        Name: ProsperRatingCombined, dtype: int64

```

We could have a potential issue with two "A" grades combining to make an "AA" grade. Let's double check the number of values before and after to make sure we didn't erroneously create "AA" grades.

```

In [91]: #Are the number of A and AA ratings the same before and after combining the columns?
(df["ProsperRatingCombined"] == "A").sum() == (df.CreditGrade == "A").sum() + (df["ProsperRating (Alpha)"] == "A").sum()
(df["ProsperRatingCombined"] == "AA").sum() == (df.CreditGrade == "AA").sum() + (df["ProsperRating (Alpha)"] == "A").sum()

Out[91]: True

```

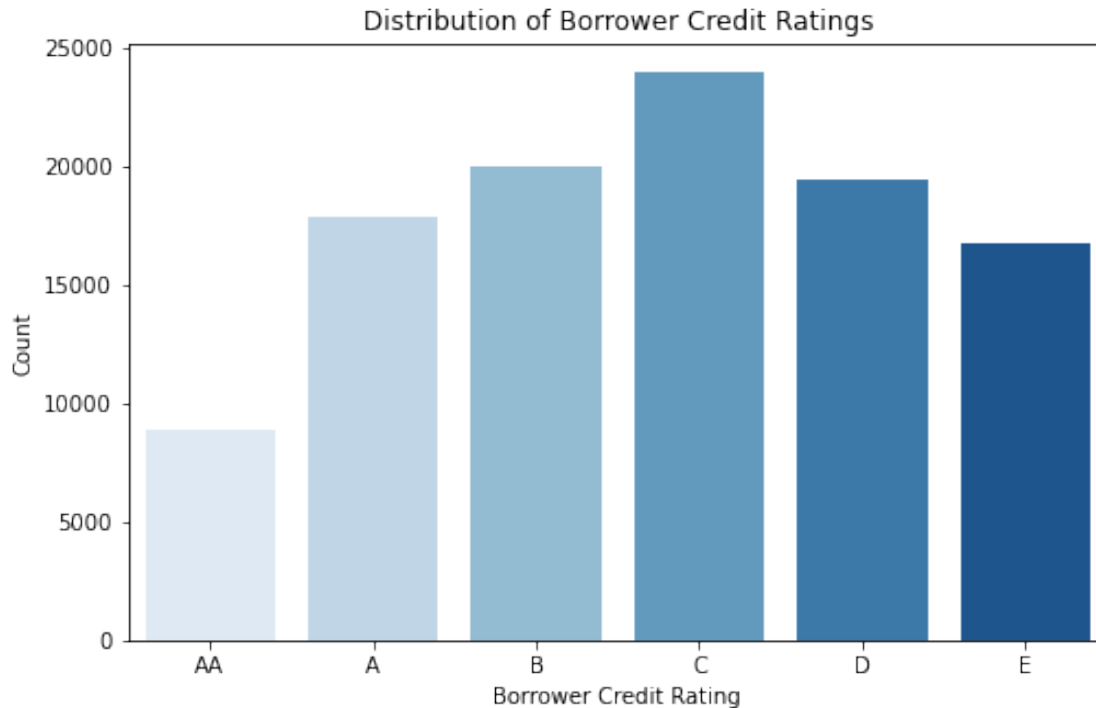
Great! We didn't introduce any erroneous data since there were no rows with an 'A' rating for both their Credit Grade and Prosper Rating.

```

In [92]: #Convert string to ordinal category type
ordinal_rating = ["AA", "A", "B", "C", "D", "E"]
ordered_var = pd.api.types.CategoricalDtype(ordered = True, categories = ordinal_rating)
df["ProsperRatingCombined"] = df["ProsperRatingCombined"].astype(ordered_var)

In [93]: plt.figure(figsize = [8, 5])
sns.countplot(data = df, x = "ProsperRatingCombined", palette = "Blues")
plt.xlabel("Borrower Credit Rating")
plt.ylabel("Count")
plt.title("Distribution of Borrower Credit Ratings");

```



The distribution of credit ratings is skewed to the right which coincides with our observation of estimated return distribution also being skewed to the right


This could mean that for a particular credit rating we have a range of possible estimated returns, let's look into this further in our bivariate visualizations.

1.3.4 Employment Status

As we want to look into the IncomeRange of borrowers, we wanted to check what their employment status is before looking into their income ranges.

```
In [94]: sns.countplot(data = target_df, x = "EmploymentStatus", color = base_color).set(title="")
plt.xticks(rotation = 90)
```

```
Out[94]: (array([0, 1, 2, 3, 4, 5, 6, 7]),
 [Text(0, 0, 'Self-employed'),
  Text(1, 0, 'Not available'),
  Text(2, 0, 'Full-time'),
  Text(3, 0, 'Other'),
  Text(4, 0, 'Employed'),
  Text(5, 0, 'Not employed'),
  Text(6, 0, 'Part-time'),
  Text(7, 0, 'Retired')])
```



Not surprisingly, the largest category of borrowers are employed and the largest sub-category of the employed status is full-time. Self-employed makes up the second largest sub-category of the employed status. The "other" category is most likely to be comprised of those falling into a "student" status.

1.3.5 Income Range

Let's continue our exploration by looking at a couple of the underwriting features for borrowers to be assigned these Credit Ratings. We'll start with taking a look at Income Ranges.

```
In [95]: df.IncomeRange.value_counts()
```

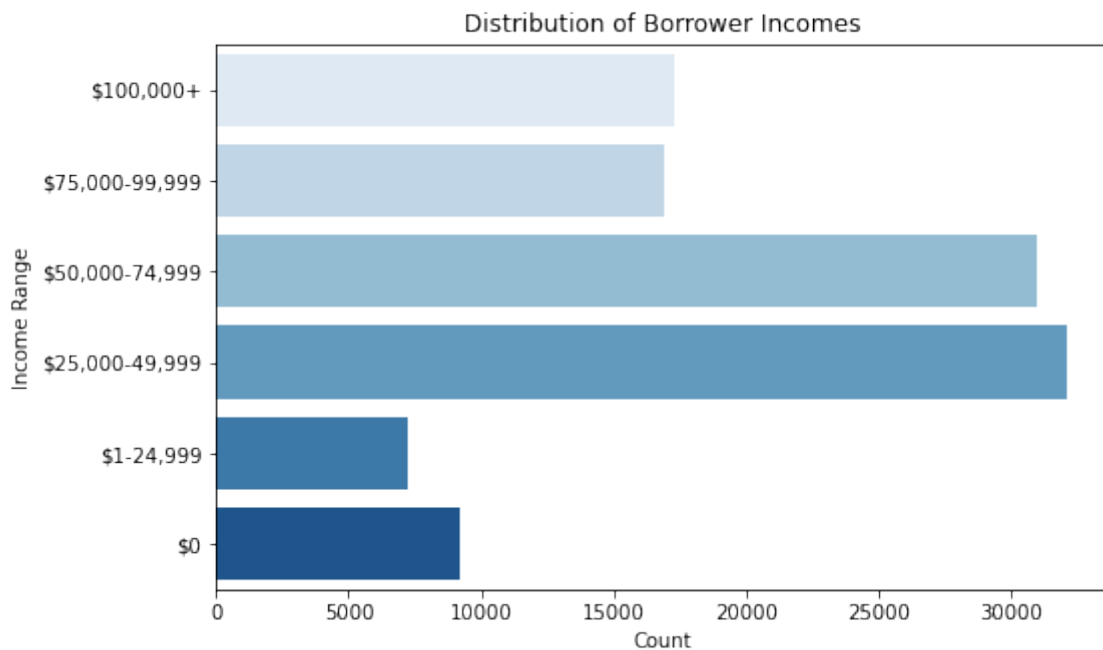
```
Out[95]: $25,000-49,999    32152
          $50,000-74,999    31005
          $100,000+         17321
          $75,000-99,999    16899
          Not displayed      7741
          $1-24,999         7261
          Not employed       806
          $0                 621
          Name: IncomeRange, dtype: int64
```


To simplify our visualization, let's combine the "Not employed", "Not displayed", and "\$0" categories together.

```
In [96]: df["IncomeRange"].replace(["Not employed", "Not displayed"], "$0", inplace = True)
```

```
In [97]: #Create ordinal categories for income
ordinal_rating = ["$100,000+", "$75,000-99,999", "$50,000-74,999", "$25,000-49,999", "$1-24,999", "$0"]
ordered_var = pd.api.types.CategoricalDtype(ordered = True, categories = ordinal_rating)
df["IncomeRange"] = df["IncomeRange"].astype(ordered_var)
```

```
In [98]: #Bar chart
plt.figure(figsize = [8, 5])
sns.countplot(data = df, y = "IncomeRange", palette = "Blues")
plt.ylabel("Income Range")
plt.xlabel("Count")
plt.title("Distribution of Borrower Incomes");
```



Its interesting to see that the number of borrowers for the 25k-49k range is the same as the 50k-74k range. The amount of borrowers in the 75k-100k range is the same as all borrowers that have an income of over 100k.


1.3.6 Debt to Income Ratio

Let's see what the borrower's Debt to Income Ratio looks like.

```
In [99]: df.DebtToIncomeRatio.describe()
```

```
Out[99]: count    105263.000000
         mean      0.275976
         std       0.551811
         min       0.000000
         25%       0.140000
         50%       0.220000
         75%       0.320000
         max       10.010000
         Name: DebtToIncomeRatio, dtype: float64
```

```
In [100]: #Distribution plot
plt.figure(figsize = [8,5])
sns.distplot(df.DebtToIncomeRatio, kde = False, bins = 200)
plt.xlim(-.1, 1)
plt.xlabel("Debt to Income Ratio")
plt.ylabel("Count")
plt.title("Distribution of Borrower Debt to Income Ratio");
```



It appears that most borrowers who apply for a loan have a debt to income ratio around 20%.

1.3.7 Discuss the distribution(s) of your variable(s) of interest: Were there any unusual points? Did you need to perform any transformations?

We saw relatively normal distributions for income, debt to income, credit ratings, and estimated returns. However, trying to calculate the real time actual returns proved to be difficult. We may

want to revisit this feature if we are trying to answer a particular question about the data set. The majority of borrowers are employed and a majority of those are full-time.

1.3.8 Of the features you investigated, were there any unusual distributions? Did you perform any operations on the data to tidy, adjust, or change the form of the data? If so, why did you do this?

We saw that from our simplified actual returns formula, that the values returned vary greatly between different loan statuses.

To tidy up our visualizations, we combined several rows for different visuals where the categories appeared to be redundant. Such as the income ranges of "Not employed" and "Not displayed" are essentially equivalent to "\$0". For Credit Ratings, Prosper developed their proprietary system after 2009 which resulted in many missing rows. We combined these features to get a complete Credit Rating feature. We also combined the "High Risk", "No Credit", and "E" categories together.

1.4 Bivariate Exploration

First let's look at the pairwise correlation between the numeric features we're interested in and see if there is any actual relationship between them.


```
In [103]: #Numeric and categorical variables
numeric_vars = ["DebtToIncomeRatio", "ActualReturns", "EstimatedReturn"]
categoric_vars = ["IncomeRange", "ProsperRatingCombined"]

In [104]: #Plot matrix with 0.5% of data points
samples = np.random.choice(df.shape[0], int(df.shape[0]*.005), replace = False)
df_samp = df.loc[samples,:]

#Pair grid
g = sns.PairGrid(data = df_samp, vars = numeric_vars, size = 3, aspect = 1.25)
g = g.map_diag(plt.hist, bins = 50)
g.map_offdiag(plt.scatter, alpha = 1/8)

#Set axis limits
g.axes[0,0].set_ylim(0,1)
g.axes[0,0].set_xlim(-.1,1)
g.axes[0,1].set_xlim(-1.1,2)
g.axes[0,2].set_xlim(0,.2)

g.fig.suptitle("Matrix Plot of Estimated Returns, Actual Returns, and Debt to Income R")
g.fig.subplots_adjust(top = .95);
```



From the univariate exploration, we have a good understanding regarding the range of values we can expect. By limiting the axis values, we are able to prevent outliers from skewing the plot. Otherwise, most of the points will appear to be grouped in a vertical line.

Surprisingly, there doesn't appear to be any correlation between the estimated returns and the actual returns we calculated. We may have to investigate further by querying select categories of loan status.

There doesn't appear to be a correlation between any of the other numeric variables either.

1.4.1 Income Range and Credit Rating vs Estimated Returns

Let's take a look at how a borrower's income range and credit rating relates to the investor's estimated returns.

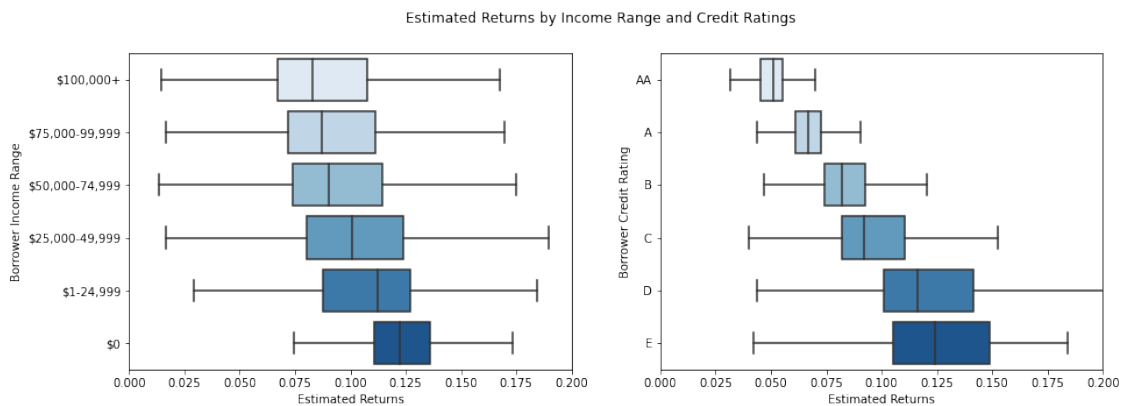
```
In [105]: #Boxplots
plt.figure(figsize = [15, 5])

plt.subplot(1, 2, 1)
base_color = sns.color_palette()[0]
sns.boxplot(data = df, y = "IncomeRange", x = "EstimatedReturn", palette = "Blues", sh
plt.xlim(0, .2)
plt.ylabel("Borrower Income Range")
```

```
plt.xlabel("Estimated Returns")

plt.subplot(1, 2, 2)
base_color = sns.color_palette()[0]
sns.boxplot(data = df, y = "ProsperRatingCombined", x = "EstimatedReturn", palette = "
plt.xlim(0,.2)
plt.ylabel("Borrower Credit Rating")
plt.xlabel("Estimated Returns")

plt.suptitle("Estimated Returns by Income Range and Credit Ratings");
```




From these boxplot, we're able to clearly see that as the borrowers income increases, the median estimated return decreases. The credit rating boxplot shows the trend more clearly, as a borrowers credit rating increases, the median estimated return greatly decreases.

Next let's plot the actual returns.

```
In [106]: #Boxplots
plt.figure(figsize = [15, 5])

plt.subplot(1, 2, 1)
base_color = sns.color_palette()[0]
sns.boxplot(data = df, y = "IncomeRange", x = "ActualReturns", palette = "Blues", show

plt.subplot(1, 2, 2)
base_color = sns.color_palette()[0]
sns.boxplot(data = df, y = "ProsperRatingCombined", x = "ActualReturns", palette = "Bl
```




This plot is quite interesting as it suggests that investing in borrowers with no reported income and the highest credit ratings will produce the highest returns.

1.4.2 Credit Start with Listing Category

In [107]: `sns.countplot(data = target_df, x = "LoanStatus", hue = "ListingCategory (numeric)", p`

Out[107]: `[Text(0.5, 1.0, 'Loan Status by Listing Category')]`



Observation 2

- In both of the graphs, the debt consolidation is the most frequent in the defaulted and completed categories


1.4.3 Credit Score VS Income Range

Let's plot the credit score along with the income range to see if there is any correlation there.

```
In [108]: #Clustered bar chart
plt.figure(figsize = [10,8])

ax = sns.countplot(data = df, y = "IncomeRange", hue = "ProsperRatingCombined",
                  palette = 'Blues')

ax.legend(title = "Borrower Credit Rating")
plt.title("Distribution of Income Ranges and Credit Ratings")
plt.ylabel("Borrower Income Range");
```



From this clustered bar chart, we can see the trend of higher incomes correlate to a higher credit rating qualifying the borrower for the best (lowest) interest rates.

Having an income of \$75k+ doesn't automatically guarantee the best credit ratings, which is curious to see. Perhaps we can see what other factors play a big role in determining one's credit rating later on.

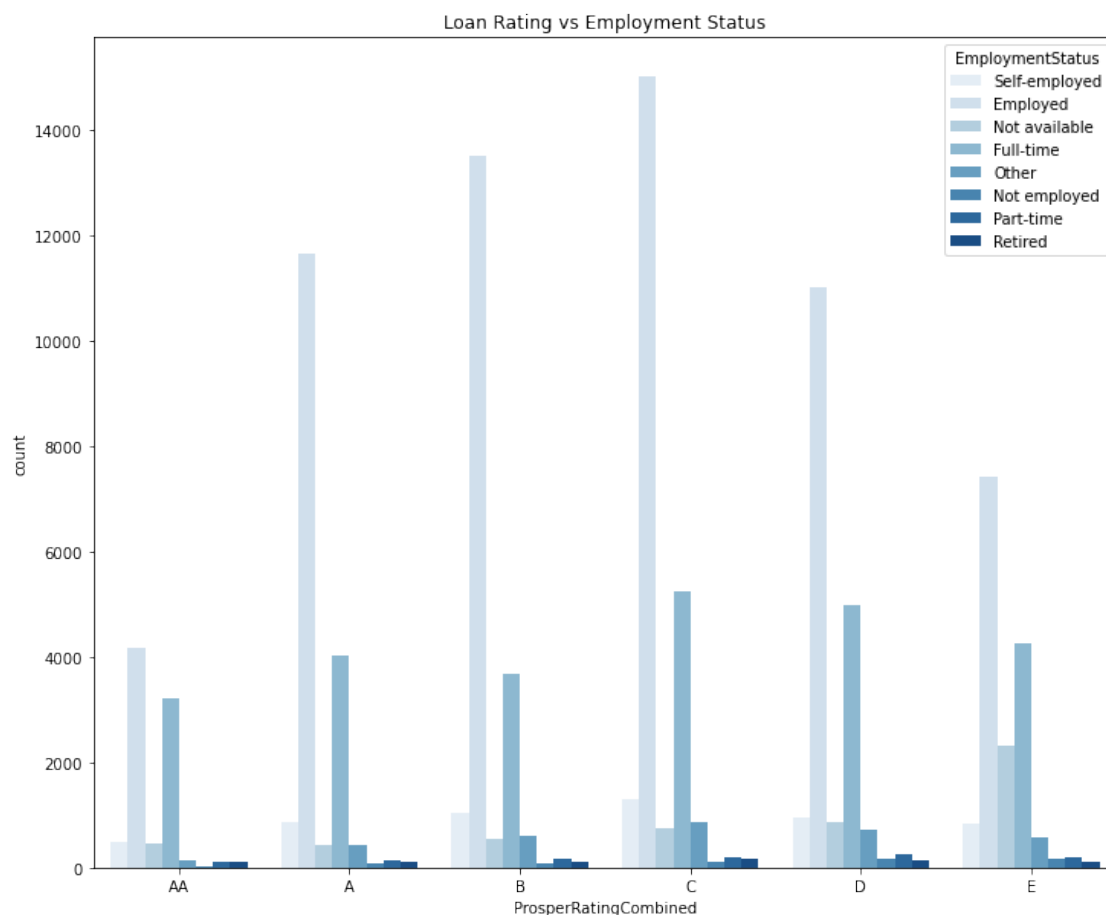
Conversely, it also appears that in the 0 income range, there are a high amount of "AA" loans relative to the 1k-25k income range.

1.4.4 Prosper Rating and Employment Status

Let's see if there is a relationship between the Prosper Rating and the Employment Status.

```
In [109]: plt.figure(figsize = [12, 10])
          sns.countplot(data = df, x = "ProsperRatingCombined", hue = "EmploymentStatus", palette="magma")
```

```
Out[109]: [Text(0.5, 1.0, 'Loan Rating vs Employment Status')]
```



Lower ratings seem to have greater proportions of individuals who selected the "Not Employed", "Self-Employed", "Retired", and "Part-Time" employment categories/sub-categories

1.5 Talk about some of the relationships you observed in this part of the investigation. How did the feature(s) of interest vary with other features in the dataset? Did you observe any interesting relationships between the other features (not the main feature(s) of interest)?

From our matrix plot, our numerical features of interest did not reveal any correlation. It was surprising to see that there was no correlation between debt to income ratios, estimated returns, and actual returns.

From our box plots, we saw a trend of median estimated returns decrease as the borrowers income range increased. This makes intuitive sense as borrowers with a higher income can demand the most competitive rates from creditors. The overlap in our boxplot suggests that borrowers in any income range (except 0) can produce the same estimated return for an investor. The trend between lower estimated returns and higher credit ratings was much more apparent with narrowing quartiles and lower medians.

The box plot for the actual returns revealed additional tweaking or grouping will be necessary to produce any insights.

From our clustered bar chart, we peeled back another layer to look at credit ratings and income ranges. This revealed that the previous trend of high income equating to better credit ratings, however there are some interesting cases where a high number of borrowers with no income qualified for the highest credit ratings. Also, we saw that borrowers with incomes of greater than 75k can still be considered "High Risk" borrowers.

EmploymentStatus of individuals with lower rated loans seemed to have selected the "Not Employed", "Self-Employed", "Retired", or "Part-Time" employment categories/sub-categories.


1.6 Multivariate Exploration

The main thing we want to explore in the multivariate exploration is the relationship between income ranges, credit scores, and estimated returns, but we may also compare other data points that come from questions pertaining to this portion of the exploration.

```
In [110]: #Create a heat map of average estimated returnsb
plt.figure(figsize = [10,5])

cat_med = df.groupby(["ProsperRatingCombined", "IncomeRange"]).median()["EstimatedReturn"]
cat_med = cat_med.reset_index(name = "EstimatedReturnMedian")
cat_med = cat_med.pivot(index = "IncomeRange", columns = "ProsperRatingCombined", values = "EstimatedReturnMedian")

sns.heatmap(cat_med, annot = True, fmt = ".1f", cmap = "Blues", cbar_kws = {"label" : "Estimated Return Median"})
plt.xlabel("Credit Rating")
plt.ylabel("Income Range")
plt.title("Heatmap of Median Estimated Returns by Credit Rating and Income Range");
```



From the heat map we continue to see the trend of lower credit ratings, and lower incomes equate to higher returns. It's interesting to see that the income range does not play a big influence on the estimated returns by credit rating category.

Across the credit ratings "AA", "A", and "B", we observe no difference in the rate of return for income ranges 1k-100k+. However, there is a definite jump in returns for the 0 income range across all credit rating categories.


Across the credit ratings "C", "D", and "E", we notice a bit more variability however the median return is still within +- 1% (except the 0 category).

1.6.1 Relationships between Credit Category, Credit Rating, and Loan Outcome

Let's see if there is a relationship between the credit category, rating, and the outcome.

```
In [111]: sns.catplot(x = "ProsperRatingCombined", hue = "LoanStatus", col = "ListingCategory (n",
                    data = target_df, kind = "count", palette = "Blues", col_wrap = 3).set(tit
```

```
Out[111]: <seaborn.axisgrid.FacetGrid at 0x242316ff6a0>
```

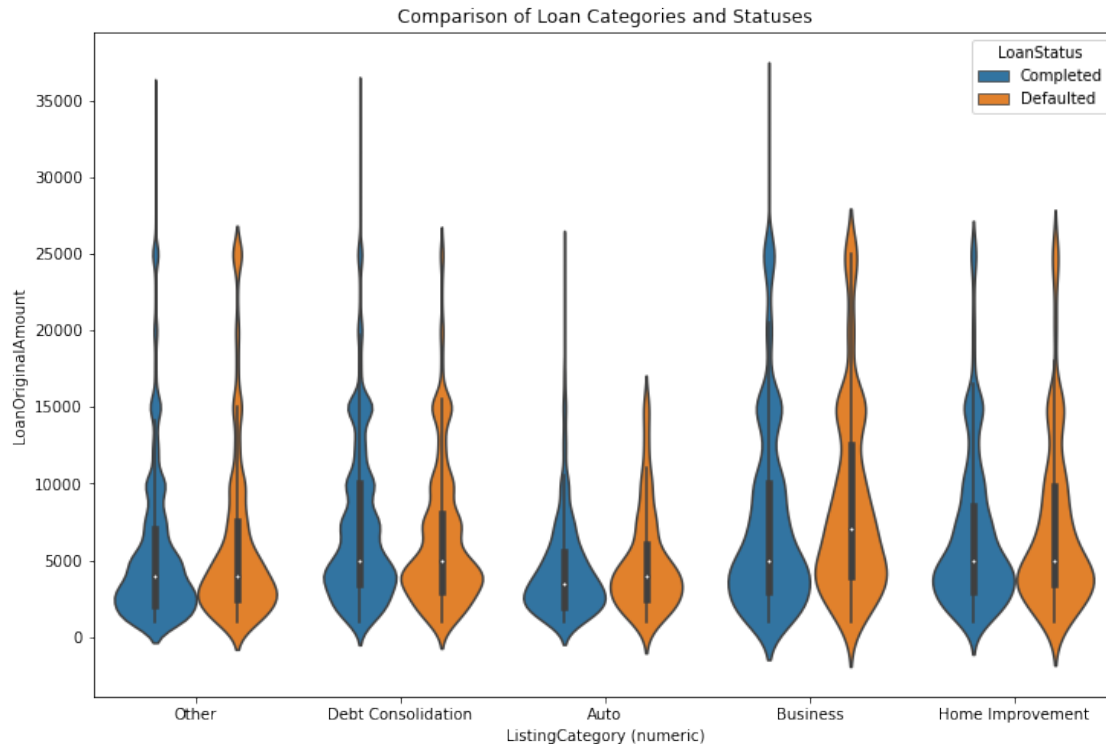


There is no substantial difference for defaulted laons in different categories when broken up by the ratings.

1.7 Amount, Listing Category Loan, and Loan Status Impact

```
In [112]: plt.figure(figsize = [12, 8])
           sns.violinplot(data = target_df, x = "ListingCategory (numeric)", y = "LoanOriginalAmount")

Out[112]: [Text(0.5, 1.0, 'Comparison of Loan Categories and Statuses')]
```



Except for Auto, the Business and Home Improvement categories don't have equal means between the loan statuses. The business category seems to have a larger amount of loans distributed and Debt Consolidation loans come in second to the Business category.

1.8 Talk about some of the relationships you observed in this part of the investigation. Were there features that strengthened each other in terms of looking at your feature(s) of interest Were there any interesting or surprising interactions between features?

Our initial assumptions were strengthened. Most of the defaulted loans comes from individuals with low Prosper ratings and the Business category had the largest amount of loans. We simplified the plots a bit by only looking at three loan statuses: Completed, Current, and Defaulted.

For loans that are defaulted, there doesn't appear to be any correlation between the predicted estimated return and the actual return of the loan.

For loans that are current, there appears to be a linear relationship between the estimated and actual returns! Recall that during our bivariate exploration we were unable to see the trend as all of the loan status were lumped together.

It was interesting to find that defaulted loans for individuals with high Prosper ratings tended to be larger than completed loans.